

# Introduction to Dplyr

Federal Reserve Board of Governors  
Howard University

# Introduction

- ▶ Last week - visualization with ggplot2
- ▶ Today - Dplyr
  - ▶ Data Analysis
  - ▶ Data Manipulation
- ▶ More background on this topic can be found in *R 4 Data Science* chapter 5 <http://r4ds.had.co.nz/transform.html>

# Today's Question

- ▶ American Community Survey for the state of California
  - ▶ income, education, age, gender,
  - ▶ 2010-2014
- ▶ How does income differ for men and women?
- ▶ How does income differ by age?

# Data

- ▶ `ca_acs.csv`
  - ▶ Read in this file as a `data.frame` and call it: “acs”
- ▶ Disclaimer: Data are taken from a random 20,000 subsample of the 5-year ACS Public Use data for California; results should be considered only as instructional exercises

# Examining our Data

- ▶ How many columns?
- ▶ Names of our columns?
- ▶ Data is messy
- ▶ Install and attach the dplyr package

# Guiding Question

- ▶ Easier to analyze data if you have a question/goal
- ▶ How do wages and income differ for men and women in California by age?
  - ▶ Line plot with age on the x-axis and dollars on the y-axis
- ▶ First step?

# Analyzing our columns

- ▶ names are cryptic acronyms
  - ▶ Data dictionaries
  - ▶ Census provides a data dictionary for the ACS at [www2.census.gov/programs-surveys/acs/tech\\_docs/pums/data\\_dict/PUMS\\_Data\\_Dictionary\\_2011-2015.pdf](http://www2.census.gov/programs-surveys/acs/tech_docs/pums/data_dict/PUMS_Data_Dictionary_2011-2015.pdf)
- ▶ How do wages and income differ for men and women in California by age

## Selecting our Columns

- ▶ Age (AGEP), Gender (SEX), wages (WAGP), income (PINCP)
  - ▶ inflation adjustment factor (ADJINC)
  - ▶ weighting factor (PWGTP)



# dplyr::select

- ▶ `Select()` allows us to subset columns from `data.frames`
  - ▶ `select(.data, ...)`
  - ▶ `.data = data.frame`
  - ▶ `... = columns`
- ▶ `Select()` example

## Deselecting Columns

- ▶ Putting the minus sign, -, in front of a column name deselects it
  - ▶ As output you would get every column except for the one you deselect

```
names(income_data)
```

```
Output: [1] "AGEP"    "SEX"     "WAGP"    "WKHP"    "PINCP"
```

```
Output: [6] "ADJINC"  "PWGTP"   "PUMA10"
```

```
deselected_data <- select(income_data, -PUMA10, -AGEP)  
names(deselected_data)
```

```
Output: [1] "SEX"     "WAGP"    "WKHP"    "PINCP"   "ADJINC"
```

```
Output: [6] "PWGTP"
```

## In class exercise: Select

- ▶ Now it's your turn, create a data.frame called `exercise_1_data` of the following columns from `acs`: `SEX`, `WKHP`, `ADJINC`, `PWGTP`, `WAGP`

# Filtering our data

- ▶ Remove unusable observations
  - ▶ rows with NA values in them
- ▶ `filter()`
  - ▶ `filter(.data, ...)`
    - ▶ `.data = data.frame`
    - ▶ `... = logical conditions (>, <, <=, >=, etc...)`
- ▶ `Filter()` example 1

# Boolean Logic

- ▶ Comparison statements that evaluate to TRUE or FALSE
  - ▶ string multiple together with & (and) and | (or) operators
  - ▶  $3 < 4$  evaluates to TRUE,
  - ▶  $3 < 4 \& 5 > 7$  evaluates to FALSE
    - ▶ When making a statement with & remember: TRUE & TRUE = TRUE while any other combination = FALSE
  - ▶ However, the statement  $3 < 4 \mid 5 > 7$  evaluates to TRUE
    - ▶ Only one statement must be TRUE when using | for the whole statement to be TRUE

## Boolean logic: exercises

- ▶ Remember that `==` is “equal to” in R, since `=` also has assignment powers in R, like `<-`
- ▶ `%in%` is the “in” operator, for example, `2 %in% c(1, 2, 3)` is `TRUE`
  - ▶ `filter(acs, AGE < 30)`
  - ▶ `filter(acs, AGE < 30 & SEX == 1)`
  - ▶ `filter(acs, AGE < 25 | AGE > 65)`
  - ▶ `filter(acs, Occ %in% c("Computer/Math", "Legal"))`

## Back to Filter

- ▶ `filter()` can take multiple criteria
  - ▶ Example above use |
- ▶ Adding more conditions with `,` is the same as `&`

```
identical(filter(income_data, AGE == 24 & SEX == 1),  
          filter(income_data, AGE == 24, SEX == 1))
```

Output: [1] TRUE

What do you think happens if you do not provide a search criterion to `filter()`?

## In class exercise: Filter

- ▶ Re-assign your exercise\_1\_data data.frame to exercise\_2\_data and filter out any NA values in the columns
- ▶ Additionally, make sure that your WKHP variable is greater or equal to 10 and less than or equal to 60
- ▶ Your data should have the following dimensions:

8598, 5



# Creating a Workflow

- ▶ Need to execute multiple steps
- ▶ One method: create new objects

```
object1 <- select(acs, AGE, SEX, WAGE,
                  PINCP, ADJINC, PWGTP)

object2 <- filter(object1,
                  !(is.na(AGE) | is.na(PINCP)))
```

# Combining statements

- ▶ Creating new objects each time
  - ▶ Time consuming and confusing
- ▶ Nesting functions
  - ▶  $f(g(x)) = 5$

## Nested Functions example

```
income_data_nest <- filter(select(acs, AGE, SEX, WAGE,  
                                PINCP, ADJINC, PWGTP),  
                           !(is.na(AGE) | is.na(PINCP)))  
  
identical(object2, income_data_nest)
```

Output: [1] TRUE

Also confusing!

# Pipe Operators

- ▶ `%>%` operator from dplyr
  - ▶ output of the left side becomes input on the right side
  - ▶  $x \%>\% f() == f(x)$

```
income_data_pipe <- acs %>%  
  select(AGEP, SEX, WAGP, PINCP, ADJINC, PWGTP) %>%  
  filter(!(is.na(AGEP) | is.na(PINCP)))  
  
identical(income_data_nest, income_data_pipe)
```

Output: [1] TRUE

## Pipes Explained

- ▶ By default the pipe operator takes the output from the left function and passes it to the right function **as the first argument**
  - ▶ Can change manually with the `.` operator

```
25 %>% seq(30, by = 1)
```

```
Output: [1] 25 26 27 28 29 30
```

```
25 %>% seq(30, ., by = -1)
```

```
Output: [1] 30 29 28 27 26 25
```

- ▶ The first example is `seq(25, 30, by = 1)` while the second is `seq(30, 25, by = -1)`

# Mutating our data

- ▶ Adjust our income and wage data to be in constant 2014 USD
  - ▶ Move the decimal in ADJINC over
  - ▶ Recode our gender column to be “Male” and “Female”
  - ▶ Create buckets for our weekly hours column
- ▶ `mutate()`
  - ▶ calculate new columns, overwrite current columns, or delete columns
- ▶ `mutate()` example
  - ▶ Can combine mutations into a single `mutate()` call

## Updating Income and Wages

- ▶ Can refer to a column created earlier in a `mutate()` call within that same function call!
- ▶ `mutate()` example 2

# Updating Gender

- ▶ Use a function within `mutate()`
  - ▶ `ifelse()`
  - ▶ `plyr::round_any()`
    - ▶ Need `plyr` installed
    - ▶ `package::function` notation

Output:	SEX	gender	WKHP	Hours
Output: 1	2	Female	40	40
Output: 2	1	Male	40	40
Output: 3	2	Female	40	40
Output: 4	1	Male	60	60
Output: 5	2	Female	30	30
Output: 6	2	Female	45	45



## case\_when() a close cousin of ifelse

- ▶ Ifelse is good when there are two options (i.e. male or female)
- ▶ What do you do when you have many different possibilities?

```
adjusted_data %>%  
  mutate(gender = case_when(SEX == 2 ~ "Female",  
                             SEX == 1 ~ "Male")) %>%  
  select(SEX, gender) %>%  
  head(3)
```

Output:   SEX gender

Output: 1   2 Female

Output: 2   1   Male

Output: 3   2 Female

## In class Exercise: Mutate

- ▶ Now it's time to update your data: `exercise_2_data`
- ▶ Update `SEX`, `ADJINC`, and `WKHP` as seen previously
- ▶ Update `WAGP` to account for the adjustment factor

Output:	SEX	WKHP	ADJINC	PWGTP	WAGP	Gender	Hours
Output: 1	2	40	1.024037	15	23552.85	Female	40
Output: 2	1	40	1.024037	22	15360.56	Male	40
Output: 3	2	40	1.024037	7	38913.41	Female	40
Output: 4	1	60	1.008425	53	25412.31	Male	60
Output: 5	2	30	1.094136	10	5470.68	Female	30
Output: 6	2	45	1.008425	19	93783.52	Female	45

## Deleting your columns

- ▶ using `mutate()` set your column = NULL

```
"SEX" %in% names(gender_hours_data)
```

Output: [1] TRUE

```
gender_hours_data <- gender_hours_data %>%  
  mutate(SEX = NULL,  
         WKHP = NULL)
```

```
"SEX" %in% names(gender_hours_data)
```

Output: [1] FALSE

What is the difference between this method and deselecting data?

## The fun part: Summarize

- ▶ Generate our summary statistics
- ▶ Calculate functions over entire column(s) in a data.frame.
  - ▶ mean, median, etc. . .
- ▶ Summarise example

```
Output:      income      wages
```

```
Output: 1 54836.58 48077.21
```

# Grouping

- ▶ Summarize by itself isn't all that useful
- ▶ Income for each gender/age combination.
- ▶ `group_by()`
  - ▶ When used in conjunction with summarize or other functions allows us to calculate statistics within different groups
- ▶ `ungroup()` resets our data

```
income_data_summary <- gender_hours_data %>%  
  group_by(AGEP, gender) %>%  
  dplyr::summarise(income = weighted.mean(PINCP,  
                                           PWGTP),  
                  wages = weighted.mean(WAGP, PWGTP))  
head(income_data_summary, 3)
```

Output: # A tibble: 3 x 4

Output: # Groups: AGEP [2]

Output: AGEP gender income wages

Output: <int> <chr> <dbl> <dbl>

Output: 1 20 Female 10694.49 9342.122

Output: 2 20 Male 10492.21 10178.223

Output: 3 21 Female 11375.18 11053.704

dplyr::summarise()? weigthed.mean() instead of mean()?

## In class exercise: `group_by` and `summarize`

- ▶ Using the `group_by()` and `summarize()` functions find the mean values for each gender/hour combination
- ▶ save your output as `exercise_4_data`
- ▶ Your data should look like the following:

```
Output: # A tibble: 6 x 3
```

```
Output: # Groups:   Gender [1]
```

```
Output:   Gender Hours      wages
```

```
Output:   <chr> <dbl>      <dbl>
```

```
Output: 1 Female     10  4640.393
```

```
Output: 2 Female     15 10555.757
```

```
Output: 3 Female     20 11856.598
```

```
Output: 4 Female     25 16512.445
```

```
Output: 5 Female     30 21234.293
```

```
Output: 6 Female     35 25792.914
```

# Finding the top and bottom with arrange

- ▶ sorts rows
  - ▶ The default method is ascending sorting
  - ▶ descending is possible with the `desc()` function
- ▶ `arrange()` examples



## In Class Exercise: arrange

- ▶ Using your `exercise_4_data` find the bottom five observations for average wages
  - ▶ You do not have to save your output to a variable
  - ▶ what is their gender, the hours worked?

```
Output: # A tibble: 5 x 3
```

```
Output: # Groups:   Gender [2]
```

```
Output:   Gender Hours      wages
```

```
Output:   <chr> <dbl>      <dbl>
```

```
Output: 1 Female     10  4640.393
```

```
Output: 2   Male     10  5108.755
```

```
Output: 3   Male     15  8083.622
```

```
Output: 4   Male     20 10326.199
```


```
Output: 5 Female     15 10555.757
```

# Reshaping our Data

- ▶ Long data is easier to use with ggplot2
- ▶ `Tidyr::gather()` and `Tidyr::spread()`
  - ▶ `gather` takes multiple columns and collapses into key/value pairs (wide -> long)
    - ▶ duplicates other data as needed
  - ▶ `spread` takes key-value pairs and spreads the data across multiple columns (long -> wide)
- ▶ `gather()` example

# Answering our Question

- Use ggplot to analyze our output



# Understanding your analysis

- ▶ Difference between income and wages?
- ▶ Men and women?
  - ▶ What do we not take into account?
- ▶ Jagged peaks and valleys?
- ▶ Is this accurate?


# Widening our age ranges



What is the message of this chart compared with our first chart using only 1-year age buckets?

## In Class Exercise

- Using the `exercise_3_data` create a chart showing mean wages for men and women by hours worked



What message does this chart show about wage inequality compared with our previous chart? Any other questions?

# Joins

- ▶ Need multiple datasets to talk to each other

## Combine Variables

x			y		
A	B	C	A	B	D
a	t	1	a	t	3
b	u	2	b	u	2
c	v	3	d	w	1

+

=

Use **bind\_cols()** to paste tables beside each other as they are.

A	B	C	A	B	D
a	t	1	a	t	3
b	u	2	b	u	2
c	v	3	d	w	1

**bind\_cols(...)**

Returns tables placed side by side as a single table.

BE SURE THAT ROWS ALIGN.

Use a **"Mutating Join"** to join one table to columns from another, matching values with the rows that they correspond to. Each join retains a different combination of values from the tables.

A	B	C	D
a	t	1	3
b	u	2	2
c	v	3	NA

**left\_join(x, y, by = NULL, copy = FALSE, suffix = c("x", "y"), ...)**  
Join matching values from y to x.

A	B	C	D
a	t	1	3
b	u	2	2
d	w	NA	1

**right\_join(x, y, by = NULL, copy = FALSE, suffix = c("x", "y"), ...)**  
Join matching values from x to y.

A	B	C	D
a	t	1	3
b	u	2	2

**inner\_join(x, y, by = NULL, copy = FALSE, suffix = c("x", "y"), ...)**  
Join data. Retain only rows with matches.

A	B	C	D
a	t	1	3
b	u	2	2
c	v	3	NA
d	w	NA	1

**full\_join(x, y, by = NULL, copy = FALSE, suffix = c("x", "y"), ...)**  
Join data. Retain all values, all rows.

## New Questions

- ▶ In what county do men have the highest wages? Women?
- ▶ In what county do men and women have the largest income ratio?
- ▶ Need to add county level data to our acs data
  - ▶ puma10\_county\_xwalk.csv

Output:    puma12    county

Output: 1       101 Alameda

Output: 2       102 Alameda

Output: 3       103 Alameda



## Joining our Data

- ▶ We want to use an inner join

```
Output: [1] "AGEP"    "WAGP"    "PINCP"   "ADJINC"  "PWGTP"
```

```
Output: [6] "PUMA10" "gender"  "Hours"   "county"
```

## County Level Statistics

- Which county has the highest male wages?

Output: # A tibble: 6 x 3

Output:	county	Female	Male
Output: *	<chr>	<dbl>	<dbl>
Output: 1	Marin	114325.96	69574.81
Output: 2	Placer	70621.52	NA
Output: 3	San Francisco	52621.21	98917.28
Output: 4	San Mateo	58002.22	93219.94
Output: 5	Santa Clara	56311.29	77299.68
Output: 6	Ventura	NA	70940.12

## County Income Ratio

- Male/Female income ratio (male wages/female wages)


```
Output: # A tibble: 6 x 2
Output: # Groups:   county [6]
Output:           county    ratio
Output:           <chr>    <dbl>
Output: 1      Imperial  2.644490
Output: 2 San Luis Obispo  2.471640
Output: 3      Ventura   2.376310
Output: 4      Merced    2.292796
Output: 5        Yolo    2.120836
Output: 6        Lake    1.967070
```

Let's now go back to how hours worked impacts wages

## Challenge exercise 1

- ▶ Wages seem to increase dramatically with hours worked
  - ▶ Next step: bar chart showing the percentage of men and women in each 5 hour bucket
1. Find the total number of observations in each hour/gender combination
    - ▶ `group_by()` and `dplyr::summarise()`
  2. Change your grouping to just the hour buckets
    - ▶ use `mutate` to calculate the total number of observations in each bucket and the the percentage of men/women in each bucket
  3. Plot with `ggplot()` and `geom_bar()` etc..., your y value should be the percentage
    - ▶ in the `geom_bar()` make sure to include `stat = "identity"`

## Challenge Exercise 1 Answer




Data from ACS

New insights? Further questions?

## Challenge Exercise 2

- ▶ Using dplyr and ggplot analyze how wage (WAGP) changes over age between men and women
  - ▶ control for number of hours worked
    - ▶ 10 hour buckets instead of 5
    - ▶ Limit to people who work between 20 and 60 hours per week
  - ▶ Bucket sizes for age should be 5 years
- ▶ PWGTP
- ▶ You will need to start with the original ACS data to answer this

## Challenge Exercise 2: Answer



Insights? Further Questions? Is this a “good” chart?